

Team Members:

Aysha Mohamed

Mostafa Hasona

Youssef Muhammed Abd El-alim

Emojis as Anchors to Detect Arabic Offensive Language and Hate Speech

1. Abstract

This paper introduces an innovative method for detecting offensive language and hate speech in Arabic social media posts by utilizing emojis as extralinguistic anchors. The research addresses the challenge of identifying such content across the diverse linguistic landscape of the Arabic language, which includes numerous dialects and cultural expressions. By focusing on emojis, the study circumvents the limitations of keyword-based methods that often fail to capture the full spectrum of offensive language. The collected dataset, the largest of its kind, was annotated and categorized into offensive, hate, vulgar, and violent speech. Various machine learning models were benchmarked, demonstrating superior performance compared to traditional methods. The findings underscore the importance of cultural context in detecting offensive content and provide a valuable resource for further research.

2. Introduction

This project focuses on detecting hate speech in Arabic tweets using machine learning models. The primary goals include:

- Preprocessing Arabic text data.
- Utilizing pre-trained models like BERT and fastText for classification.
- Evaluating model performance.
- Exploring adversarial robustness by perturbing input data.

Key libraries and their roles:

- **PyTorch:** A deep learning framework used for loading and running the BERT model.
- **Transformers:** A library by Hugging Face for accessing state-of-the-art pre-trained models like BERT.
- **Pandas and NumPy:** Essential for data manipulation and numerical computations.
- **NLTK:** Provides tools for tokenization and stopwords.
- **BeautifulSoup:** Used for parsing and cleaning HTML content from tweets.
- **fastText:** Used for obtaining word embeddings.

- **FarasaStemmer**: Specifically designed for stemming Arabic words.
- **emoji**: For handling and processing emojis present in the text.

3. Related Work

Previous studies have primarily focused on English-language datasets for hate speech detection, with fewer resources available for Arabic. Notable work includes the use of traditional machine learning methods and more recent deep learning approaches like BERT. However, the robustness of these models against adversarial attacks remains underexplored. Our research extends this literature by specifically targeting Arabic language challenges and assessing model resilience to adversarial perturbations, a gap not fully addressed in existing studies

4. Methods

Files and URLs:

- **Tweets File**: OSACT2022-sharedTask-test-tweets.txt containing tweet texts.
- **Labels File**: OSACT2022-sharedTask-test-taskB-gold-labels.txt containing the corresponding hate speech labels.
- **Arabic Stop Words**: Arabic_stop_words.txt lists common Arabic stop words.
- **Pre-trained Model Files**: qarib pretrained model and fasttext-ar-vectors-150.bin provide the necessary pre-trained models.

Data Preparation and Pre-processing

Train data: `"https://alt.qcri.org/resources/OSACT2022/OSACT2022-sharedTask-train.txt"`

Dev data: `"https://alt.qcri.org/resources/OSACT2022/OSACT2022-sharedTask-dev.txt"`

Test data: `"https://alt.qcri.org/resources/OSACT2022/OSACT2022-sharedTask-test-tweets.txt"`

labels for test: `"https://alt.qcri.org/resources1/OSACT2022/OSACT2022-sharedTask-test-taskB-gold-labels.txt"`

1- Data Loading

The tweet data and labels are loaded into a pandas DataFrame

Columns are renamed for consistency and clarity.

The data includes:

- **tweet**: The tweet text.
- **class**: The label indicating whether the tweet is hate speech (HS) or not (NOT_HS).

```
train_data = pd.read_csv("OSACT2022-sharedTask-train.txt", sep="\t", quoting=csv.QUOTE_NONE)
dev_data = pd.read_csv("OSACT2022-sharedTask-dev.txt", sep="\t", quoting=csv.QUOTE_NONE)
test_data = pd.read_csv("OSACT2022-sharedTask-test-tweets.txt", sep="\t", quoting=csv.QUOTE_NONE)
```

train_data	1	@USER ردينا ع القنطرة 🇲🇪	OFF	NOT_HS	NOT_VLG	NOT_VIO	
0	2	وصارت فطير البقالات عطاء صبي 🇲🇪 URL	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
1	3	@USER زوجي لو بده تلقى الشاه كتير بن ماحد زينه...	OFF	NOT_HS	NOT_VLG	NOT_VIO	
2	4	@USER مش حارب بقي <LF> مش بين حاجة خالص	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
3	5	يقوله: 90% من المسلمي <LF> <LF> اليوم الاتي...	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
4	6	حدثنا ماخبطها في في اساسا 🇲🇪 URL	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
...	
8881	8883	@USER الله يلعنهم 🇲🇪	OFF	NOT_HS	NOT_VLG	NOT_VIO	
8882	8884	حبيتي كم <LF> اناك سل زوجة بعد كم سنة زواج	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
8883	8885	@USER بالله روح زي الشعير واخلف الشو الي سويته	OFF	NOT_HS	NOT_VLG	NOT_VIO	
8884	8886	لما الشكاب تنهو هو عليك 🇲🇪 وانت التي	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	
8885	8887	... الله يلعنكم في بنات خلوني <LF> <LF> بايع الكلب	NOT_OFF	NOT_HS	NOT_VLG	NOT_VIO	

8886 rows x 6 columns

2- Text Cleaning and Normalization

Several functions are defined to preprocess the tweets:

1.Normalization:

- Replace variations of certain Arabic characters with a standard form.
- Remove Arabic diacritics which can change the meaning of words.

2.Tokenization and Stop Words Removal:

- Tokenize the text using NLTK's WordPunctTokenizer.
- Remove English and Arabic stop words using lists defined in NLTK and Arabic_stop_words

3.Additional Cleaning:

- Remove repeating characters (e.g., "ههههه" becomes "هه").
- Remove punctuation marks using a custom list of Arabic and English punctuation.

4.HTML and Non-Arabic Content Removal:

- Clean HTML tags using BeautifulSoup.
- Remove non-Arabic content and URLs.

5.Function Integration:

- The clean_str function integrates all these steps, providing a final cleaned version of the tweet text.

Train_data after preprocessing:

Dev_data after preprocessing:

Test_data after preprocessing:

```
test_data = test_data.join(test_labels)
test_data
```



	Text	label
0	مشفتش العرض بتاعهم لا مش مهمة لا URL	NOT_HS
1	RT @USER عندما تكون لوحده تحس انك لحدك صح <L...	NOT_HS
2	RT @USER ماشاء الله الرجال باين عليه محترم <LF...	NOT_HS
3	@USER شسالة احد يفهمني	NOT_HS
4	@USER افوووول استريح عاد احتفالاتنا تحط اغاني	NOT_HS
...
2535	... فله حيا وين اهلهم ذوي الله لايبلانا لهادرجه	NOT_HS
2536	RT @USER ثم الطحليه URL	NOT_HS
2537	يا وجه الله من اليوم ورايح شاورما انسى URL	NOT_HS
2538	@USER متخلف حتى الحلال حرمتوه	NOT_HS
2539	@USER حنا حقينا على بنت رئيس مو على بياح كليجا	NOT_HS

2540 rows x 2 columns

3. Data Augmentation

- This involves altering the input tweets to create new examples. The goal is to test the robustness of the model by checking if small changes in the input can change the classification.
- The process includes replacing words with similar ones using fastText embeddings and checking if the classification changes.

4. Model Setup and Loading

1.BERT-based Qarib Model

- **Model Loading:** The BERT model is loaded using the transformers library
- **Tokenizer:** The tokenizer is responsible for converting text into a format that the model can understand. It handles tasks like tokenization and padding.
- **State Dictionary:** Pre-trained weights are loaded to initialize the model with learned parameters.

```
[ ] model_name = "qarib/bert-base-qarib"
num_labels = 2 # HS & NOT_HS
config = AutoConfig.from_pretrained(model_name,num_labels=num_labels, output_attentions=True)
tokenizer = AutoTokenizer.from_pretrained(model_name,
                                         do_lower_case=False,
                                         do_basic_tokenize=True,
                                         never_split=NEVER_SPLIT_TOKENS)

tokenizer.max_len = 64
model = BertForSequenceClassification.from_pretrained(model_name, config=config)

train_dataset = SingleSentenceClassificationProcessor(mode='classification')
dev_dataset = SingleSentenceClassificationProcessor(mode='classification')

train_dataset.add_examples(texts_or_text_and_labels=train_df['text'],labels=train_df['label'],overwrite_examples = True)
dev_dataset.add_examples(texts_or_text_and_labels=dev_df['text'],labels=dev_df['label'],overwrite_examples = True)
print(train_dataset.examples[0])

train_features = train_dataset.get_features(tokenizer = tokenizer, max_length =64)
dev_features = dev_dataset.get_features(tokenizer = tokenizer, max_length =64)
# print(config)


print(len(train_features))
print(len(dev_features))
```

Training the model:

T CODE T TEXT | COPY TO FILE

```
[ ] trainer = Trainer(model=model,
                    args = training_args,
                    train_dataset = train_features,
                    eval_dataset = dev_features,
                    compute_metrics = compute_metrics)

# start training
trainer.train()
```

 BertSdpaSelfAttention is used but `torch.nn.functional.scaled_dot_product_attention` does not support non-absolute `pos` [2973/2973 11:30, Epoch 3/3]

Step	Training Loss
100	0.409300
200	0.280100
300	0.245600
400	0.208000
500	0.178200
600	0.148400
700	0.139900
800	0.128100
900	0.111900
1000	0.088700
1100	0.050600

2.fastText Model

- **Loading:** The fastText model is loaded for word embeddings, crucial for semantic similarity tasks

5. Adversarial Attack Method:

The adversarial attack involves perturbing tweets by replacing words with synonyms or similar terms, thereby attempting to mislead the text classification model. The script includes error handling to manage potential interruptions during execution.

We present a black-box greedy approach consisting of

- 1- Replacing the word with the [MASK] token
- 2- Applying Bert to find substitute words
- 3- Utilizing a pre-trained word vector model to find the most similar substitute

```
tweet_index: 0
Original tweet: بايع الكلجاء : Original label: HS.
[2024-07-27 16:18:06,621 - farasapy_logger - ERROR]: pipe brokel error code and message: [[Errno 32] Broken pipe]. reinitialize the process., This may take sometime depending on the
Perturbed tweet: بايع الكلجاء : New Label: NOT_HS.
=====
Perturbed tweet: بايع الكلجاء : New Label: NOT_HS.
=====
Tweet index: 1
Original tweet: الجادريه : Original label: HS.
Perturbed tweet: الجادريه : New Label: NOT_HS.
=====
Perturbed tweet: الجادريه : New Label: NOT_HS.
=====
Perturbed tweet: الجادريه : New Label: NOT_HS.
=====
Tweet index: 2
Original tweet: هم مايسى حماء الوهاية : Original label: HS.
Perturbed tweet: هم مايسى حماء الوهاية : New Label: NOT_HS.
=====
Tweet index: 3
Original tweet: كليلجا : Original label: HS.
Perturbed tweet: كليلجا : New Label: NOT_HS.
=====
Perturbed tweet: كليلجا : New Label: NOT_HS.
=====
Tweet index: 4
Original tweet: بايع الكلجاء وصلنا لزم نلّول الله إيلينا اعراف الحرة من زود الحياه نستحي نطع عليها يمن اي رجال مو محرم لها وهولا تصوير وقعة حتره وتربية فطسونا مكنوت : Original label: HS.
=====
Tweet index: 5
Original tweet: اتمر نطسّم الين اي سيد نتظّم حيا يا ممتو زيد هل نطرقك على أسسك بطلع ملك قحت انت وسيدك يا كلب الين ادمه عليك : Original label: HS.
=====
Tweet index: 6
Original tweet: الحزاي بول لله : Original label: HS.
Perturbed tweet: الحزاي بول لله : New Label: NOT_HS.
```

Detailed Steps:

Emoji Check:

- Ignores words containing emoji, as they are not considered for replacement.

Word Masking:

- Masks the word to be replaced and generates candidate sequences using a masked language model (unmasker_MARBERT).

Semantic Similarity:

- Computes semantic similarity between the original and candidate words using FastText embeddings.

Selection Criteria:

- Filters out candidates with the same root as the original word or those that do not fit the sentence context.

- Selects the candidate with the highest semantic similarity that changes the tweet's classification.

- **Effectiveness:** The perturbation mechanism is designed to exploit potential vulnerabilities in the text classification model. The effectiveness of perturbations depends on the model's robustness and the quality of word replacements.
- **Pre-trained Models:** Ensure that the pre-trained models and custom functions used for word replacement and classification are properly configured and accessible.

5. Evaluation and Results

5.1 Prediction and Evaluation

- **Classification:** The cleaned tweets are passed through the model, which outputs logits. The class with the highest logit is chosen as the prediction.
- **Accuracy Calculation:** Accuracy is calculated specifically for the hate speech class, providing insights into the model's performance on this critical task.

5.2 Perturbation Analysis

- **Process:** Each tweet is perturbed by replacing words with their semantically similar counterparts, as determined by the fastText embeddings.
- **Evaluation:** The impact of these perturbations is measured by the model's change in predictions, indicating the model's robustness.

BERT-based Qarib Model Training Evaluation:

trainer.evaluate()

[80/80 00:06]


(1269, 2)
(12, 1269, 12, 64, 64)
1269

	precision	recall	f1-score	support
0	0.95	0.99	0.97	1160
1	0.77	0.50	0.60	109
accuracy			0.94	1269
macro avg	0.86	0.74	0.79	1269
weighted avg	0.94	0.94	0.94	1269

[[1144 16]
[55 54]]
{'eval_loss': 0.45425209403038025,
'eval_macro_f1': 0.7866272281835168,
'eval_macro_precision': 0.8627785058977719,
'eval_macro_recall': 0.7408098702942107,
'eval_accuracy': 0.9440504334121356,
'eval_runtime': 12.8356,
'eval_samples_per_second': 98.866,
'eval_steps_per_second': 6.233,
'epoch': 3.0}

BERT-based Qarib Model Testing Evaluation:


```
print(classification_report(test_df["label"].values, outputs, target_na
```



	precision	recall	f1-score	support
NOT_HS	0.95	0.99	0.97	2269
HS	0.82	0.56	0.67	271
accuracy			0.94	2540
macro avg	0.89	0.77	0.82	2540
weighted avg	0.94	0.94	0.94	2540

6. Future Work and Improvements

- **Model Enhancement:** Future improvements could involve experimenting with more advanced models, such as transformers or ensemble methods, to enhance classification accuracy and robustness.
- **Data Augmentation:** Incorporating additional data sources, including different dialects and contexts, can improve the model's understanding and generalization.
- **Adversarial Training:** Implementing adversarial training techniques could help the model learn to resist perturbations, making it more robust to adversarial attacks

7. Conclusion

This project showcases the application of state-of-the-art NLP techniques for detecting hate speech in Arabic tweets. The comprehensive workflow includes data preparation, model loading, prediction, and robustness testing through adversarial perturbation. The results highlight the model's capabilities and areas for potential improvement.